

```
In [ ]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np

# Providing the path in drive
file_path = 'shopping_trends_updated - shopping_trends_updated.csv'

# Loading the data
df = pd.read_csv(file_path)

# Display the first few rows of the dataset
df.head()
```

Out [ ]:

	Customer ID	Age	Gender	Item Purchased	Category	Purchase Amount (USD)	Location	Size	Color	Season
0	1	55	Male	Blouse	Clothing	53	Kentucky	L	Gray	Winter
1	2	19	Male	Sweater	Clothing	64	Maine	L	Maroon	Winter
2	3	50	Male	Jeans	Clothing	73	Massachusetts	S	Maroon	Spring
3	4	21	Male	Sandals	Footwear	90	Rhode Island	M	Maroon	Spring
4	5	45	Male	Blouse	Clothing	49	Oregon	M	Turquoise	Spring

```
In [ ]: # Set the style for the plots
sns.set(style="whitegrid")

# Plot 1: Distribution of Age by Gender
plt.figure(figsize=(10, 6))
sns.histplot(data=df, x='Age', hue='Gender', multiple='stack', kde=True)
plt.title('Distribution of Age by Gender')
plt.xlabel('Age')
plt.ylabel('Count')
plt.legend(title='Gender')
plt.show()

# Plot 2: Average Purchase Amount by Category
plt.figure(figsize=(10, 6))
sns.barplot(data=df, x='Category', y='Purchase Amount (USD)', ci=None, estimator=sum)
plt.title('Average Purchase Amount by Category')
plt.xlabel('Category')
plt.ylabel('Total Purchase Amount (USD)')
plt.show()


# Plot 3: Review Ratings by Item Purchased
plt.figure(figsize=(12, 6))
sns.boxplot(data=df, x='Item Purchased', y='Review Rating')
```

```
plt.title('Review Ratings by Item Purchased')
plt.xlabel('Item Purchased')
plt.ylabel('Review Rating')
plt.xticks(rotation=45)
plt.show()

# Plot 4: Frequency of Purchases by Payment Method
plt.figure(figsize=(10, 6))
sns.countplot(data=df, x='Payment Method', hue='Frequency of Purchases')
plt.title('Frequency of Purchases by Payment Method')
plt.xlabel('Payment Method')
plt.ylabel('Count')
plt.legend(title='Frequency of Purchases')
plt.show()

# Plot 5: Seasonal Purchases by Location
plt.figure(figsize=(12, 6))
sns.countplot(data=df, x='Location', hue='Season')
plt.title('Seasonal Purchases by Location')
plt.xlabel('Location')
plt.ylabel('Count')
plt.legend(title='Season')
plt.xticks(rotation=45)
plt.show()
```

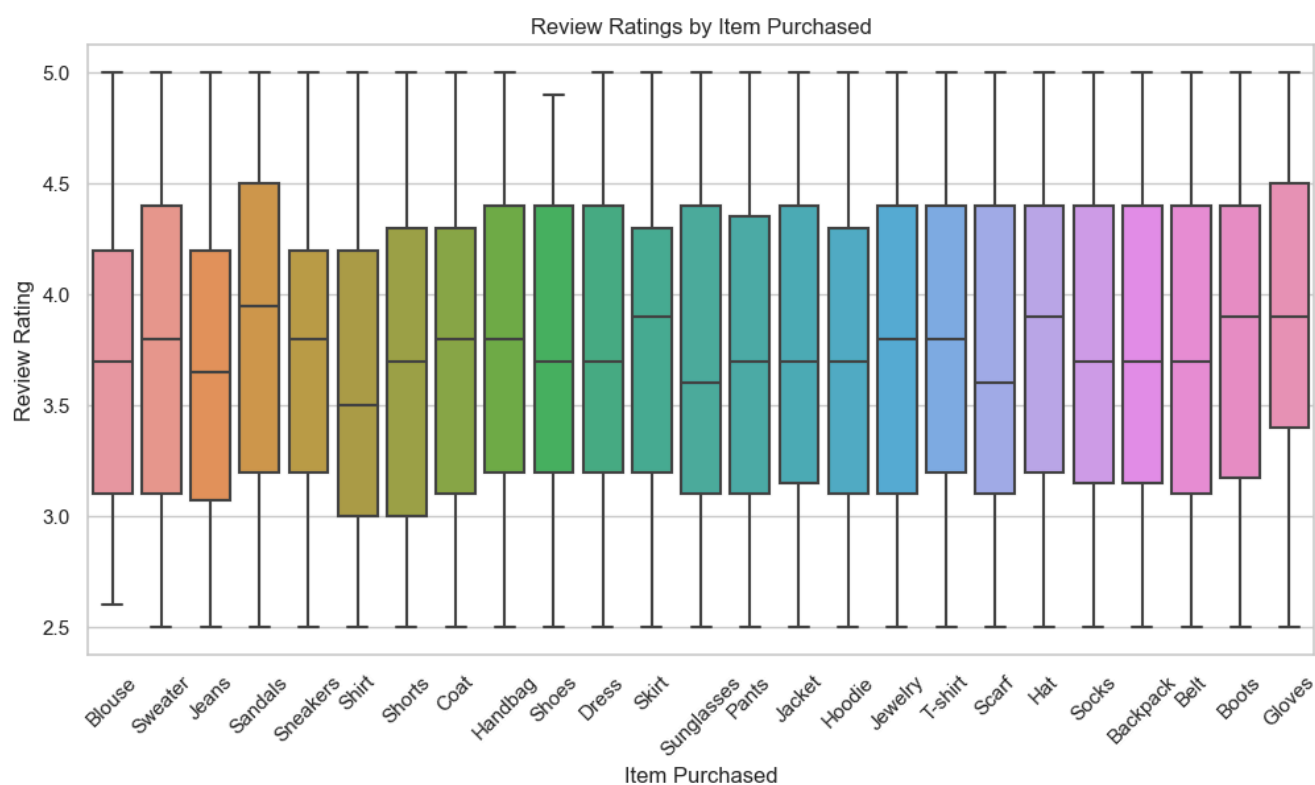
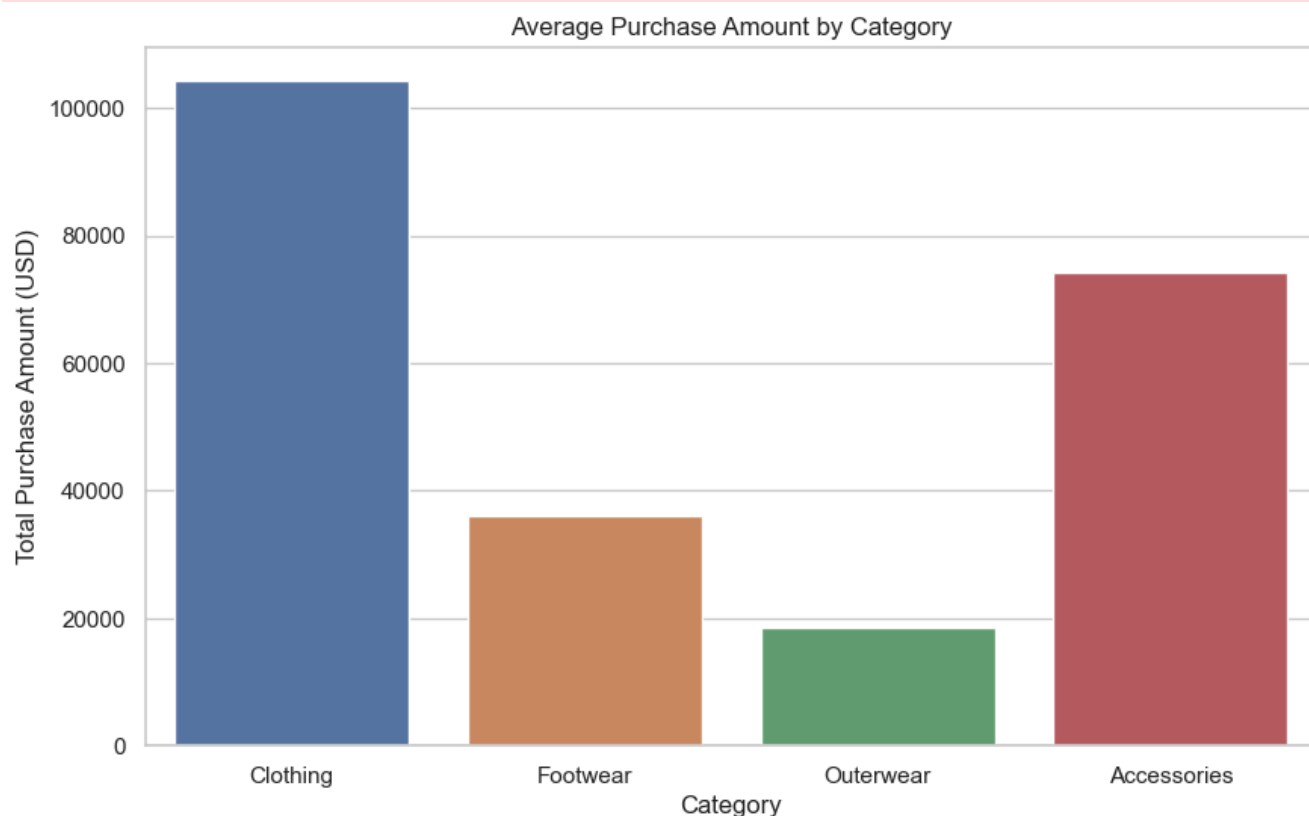
No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.

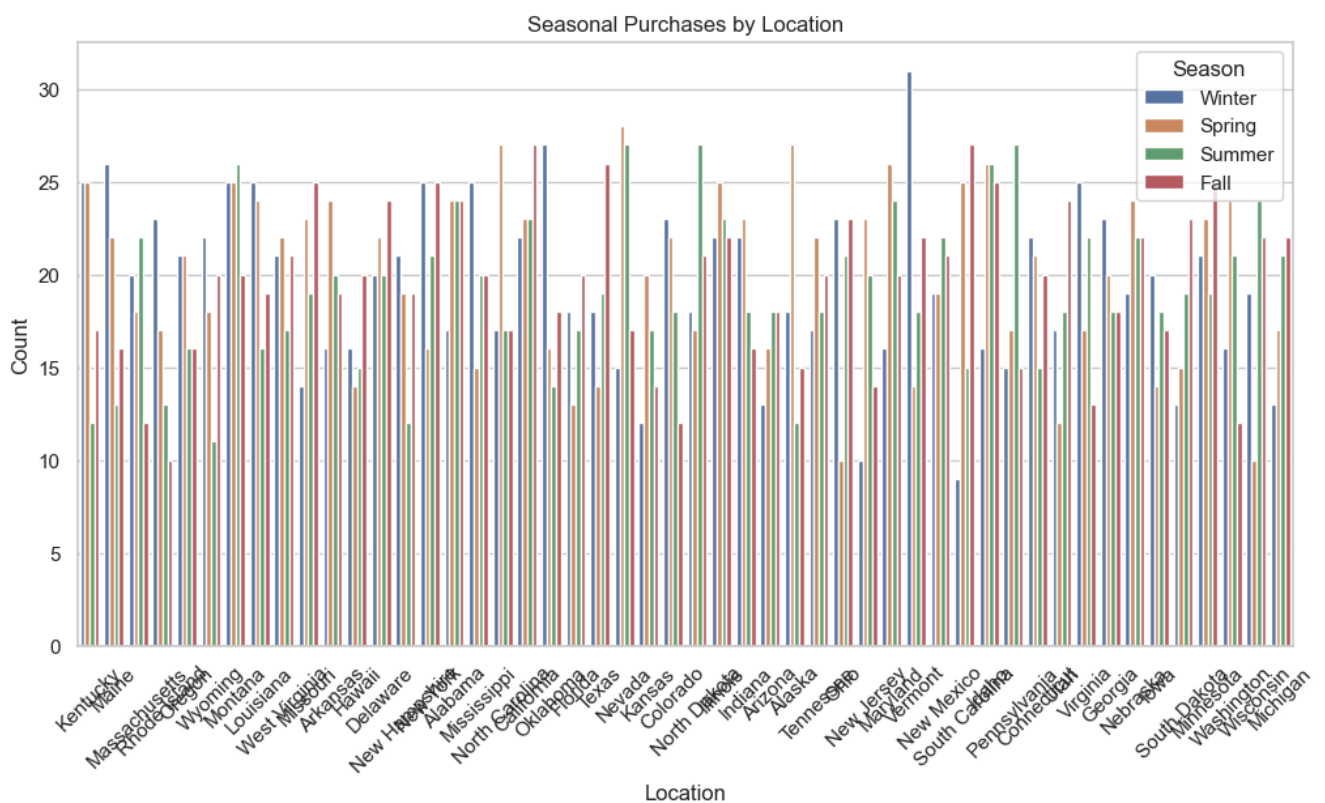



C:\Users\aravi\AppData\Local\Temp\ipykernel\_13424\3989714833.py:15: FutureWarning:

The `ci` parameter is deprecated. Use `errorbar=None` for the same effect.

```
sns.barplot(data=df, x='Category', y='Purchase Amount (USD)', ci=None, estimator=sum)
```





```
In [ ]: from sklearn.preprocessing import LabelEncoder
# Check for missing values
missing_values = df.isnull().sum()

# Encode all categorical variables using Label encoding
label_encoders = {}
for column in df.select_dtypes(include=['object']).columns:
```

```
le = LabelEncoder()
df[column] = le.fit_transform(df[column])
label_encoders[column] = le

# Check for outliers in numeric columns using IQR method
Q1 = df.quantile(0.25)
Q3 = df.quantile(0.75)
IQR = Q3 - Q1

# Define outlier criteria
outlier_criteria = (df < (Q1 - 1.5 * IQR)) | (df > (Q3 + 1.5 * IQR))

# Handle outliers: for simplicity, let's remove rows with outliers
df_cleaned = df[~outlier_criteria.any(axis=1)]

# Show the summary of the cleaned dataset
print("Missing values:\n", missing_values)
print("\nCleaned dataset info:\n")
df_cleaned.info()
print("\nPreview of the cleaned dataset:\n", df_cleaned.head())
```

Missing values:

```

Customer ID      0
Age              0
Gender           0
Item Purchased   0
Category         0
Purchase Amount (USD) 0
Location         0
Size            0
Color           0
Season          0
Review Rating    0
Subscription Status 0
Shipping Type    0
Discount Applied 0
Promo Code Used  0
Previous Purchases 0
Payment Method   0
Frequency of Purchases 0
dtype: int64

```

Cleaned dataset info:

```
<class 'pandas.core.frame.DataFrame'>
```

Index: 3576 entries, 0 to 3899

Data columns (total 18 columns):

#	Column	Non-Null Count	Dtype
0	Customer ID	3576 non-null	int64
1	Age	3576 non-null	int64
2	Gender	3576 non-null	int32
3	Item Purchased	3576 non-null	int32
4	Category	3576 non-null	int32
5	Purchase Amount (USD)	3576 non-null	int64
6	Location	3576 non-null	int32
7	Size	3576 non-null	int32
8	Color	3576 non-null	int32
9	Season	3576 non-null	int32
10	Review Rating	3576 non-null	float64
11	Subscription Status	3576 non-null	int32
12	Shipping Type	3576 non-null	int32
13	Discount Applied	3576 non-null	int32
14	Promo Code Used	3576 non-null	int32
15	Previous Purchases	3576 non-null	int64
16	Payment Method	3576 non-null	int32
17	Frequency of Purchases	3576 non-null	int32

dtypes: float64(1), int32(13), int64(4)

memory usage: 349.2 KB

Preview of the cleaned dataset:

	Customer ID	Age	Gender	Item Purchased	Category	Purchase Amount (USD)
0	1	55	1	2	1	53 \
1	2	19	1	23	1	64
2	3	50	1	11	1	73
3	4	21	1	14	2	90
4	5	45	1	2	1	49


	Location	Size	Color	Season	Review Rating	Subscription Status
0	16	0	7	3	3.1	1 \
1	18	0	12	3	3.1	1
2	20	2	12	1	3.1	1
3	38	1	12	1	3.5	1
4	36	1	21	1	2.7	1

	Shipping Type	Discount Applied	Promo Code Used	Previous Purchases
0	1	1	1	14 \
1	1	1	1	2
2	2	1	1	23
3	3	1	1	49
4	2	1	1	31

	Payment Method	Frequency of Purchases
0	5	3
1	1	3
2	2	6
3	4	6
4	4	0

```
In [ ]: # Calculate the correlation matrix
correlation_matrix = df_cleaned.corr()

# Visualize the correlation matrix using a heatmap
plt.figure(figsize=(16, 12))
sns.heatmap(correlation_matrix, annot=True, fmt=".2f", cmap="coolwarm", cbar=True)
plt.title('Correlation Matrix Heatmap')
plt.show()
```



```
In [ ]: from sklearn.metrics import r2_score
from sklearn.linear_model import LinearRegression

X = df_cleaned[['Age']]
y = df_cleaned['Purchase Amount (USD)']

model = LinearRegression()
model.fit(X, y)

# Interpret the regression coefficients
intercept = model.intercept_
slope = model.coef_[0]

print(f'Intercept: {intercept}')
print(f'Slope: {slope}')

# Determine the goodness-of-fit (R-squared value)
y_pred = model.predict(X)
r_squared = r2_score(y, y_pred)
print(f'R-squared: {r_squared}')

# Visualize the regression line on the scatter plot
plt.figure(figsize=(10, 6))
```



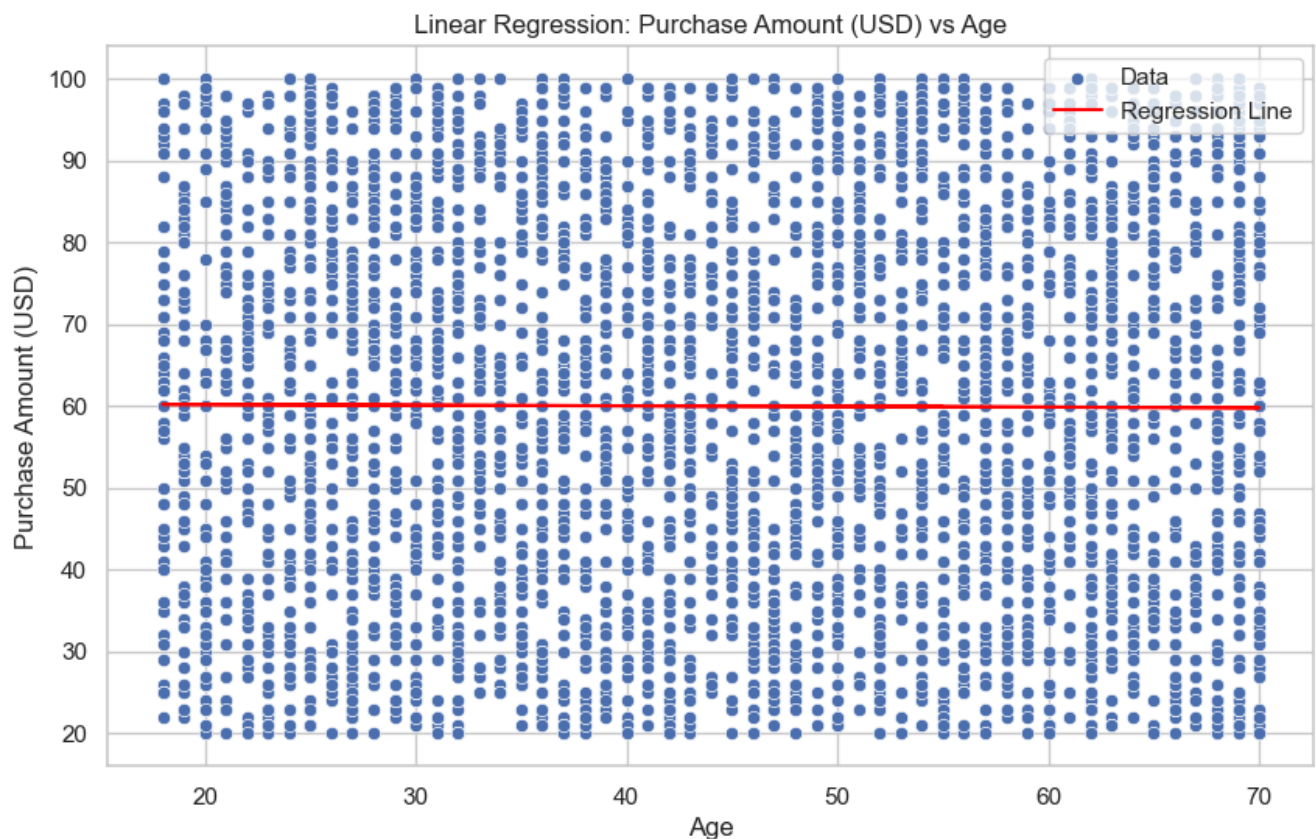
```
sns.scatterplot(x='Age', y='Purchase Amount (USD)', data=df_cleaned, label='Data')
plt.plot(X, y_pred, color='red', label='Regression Line')
plt.xlabel('Age')
plt.ylabel('Purchase Amount (USD)')
plt.title('Linear Regression: Purchase Amount (USD) vs Age')
plt.legend()
plt.show()

# Make predictions for new ages 25, 35, and 45
new_ages = np.array([[25], [35], [45]])
predictions = model.predict(new_ages)
print(f'Predicted Purchase Amounts for ages 25, 35, and 45: {predictions}')
```

Intercept: 60.35352215281485

Slope: -0.008045173122124808

R-squared: 2.6848368337106798e-05



Predicted Purchase Amounts for ages 25, 35, and 45: [60.15239282 60.07194109 59.99148936]

C:\Users\aravi\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.9\_qbz5n2kfra8p0\LocalCache\local-packages\Python39\site-packages\sklearn\base.py:493: UserWarning: X does not have valid feature names, but LinearRegression was fitted with feature names  
warnings.warn(

```
In [ ]: # • Fit a multiple linear regression model to predict Purchase Amount (USD)
# based on Age, Gender, Location, and Review Rating.
# • Interpret the regression coefficients and their significance.
# • Determine the goodness-of-fit (R-squared value).
# • Make predictions using the regression model for specific values of the
# predictors.

# Define the dependent and independent variables
X = df_cleaned[['Age', 'Gender', 'Location', 'Review Rating']]
y = df_cleaned['Purchase Amount (USD)']
```

```
# Fit the multiple linear regression model
model = LinearRegression()
model.fit(X, y)

# Get the regression coefficients and intercept
coefficients = model.coef_
intercept = model.intercept_

# Print the regression coefficients and intercept
print('Regression coefficients:', coefficients)
print('Intercept:', intercept)

# Determine the goodness-of-fit (R-squared value)
y_pred = model.predict(X)
r_squared = r2_score(y, y_pred)
print('R-squared:', r2_score(y, y_pred))

# Make predictions for specific values of the predictors
new_data = [[30, 1, 2, 4], [45, 0, 1, 3]]
predictions = model.predict(new_data)
print('Predicted purchase amounts:', predictions)

# Check the significance of the regression coefficients using a t-test
import statsmodels.api as sm

X_with_constant = sm.add_constant(X)
model_sm = sm.OLS(y, X_with_constant).fit()
print(model_sm.summary())
```

```
C:\Users\aravi\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.9_qbz5n2kfra8p0\LocalCache\local-packages\Python39\site-packages\sklearn\base.py:493: UserWarning: X does not have valid feature names, but LinearRegression was fitted with feature names
  warnings.warn(
```

Regression coefficients: [-0.00695873 -0.62055956 0.05076208 0.80223492]

Intercept: 56.48461931816324

R-squared: 0.0017198607777986208

Predicted purchase amounts: [58.96576169 58.62894331]

#### OLS Regression Results

```
=====
Dep. Variable:      Purchase Amount (USD)    R-squared:                0.002
Model:              OLS                    Adj. R-squared:           0.001
Method:             Least Squares          F-statistic:             1.538
Date:               Mon, 24 Jun 2024        Prob (F-statistic):       0.188
Time:               16:00:03                Log-Likelihood:          -16374.
No. Observations:   3576                   AIC:                     3.276e+04
Df Residuals:       3571                   BIC:                     3.279e+04
Df Model:           4
Covariance Type:    nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	56.4846	2.570	21.978	0.000	51.446	61.524
Age	-0.0070	0.026	-0.268	0.789	-0.058	0.044
Gender	-0.6206	0.845	-0.734	0.463	-2.277	1.036
Location	0.0508	0.027	1.849	0.065	-0.003	0.105
Review Rating	0.8022	0.550	1.459	0.145	-0.276	1.881

```
=====
Omnibus:            3673.807    Durbin-Watson:           1.944
Prob(Omnibus):      0.000    Jarque-Bera (JB):         222.359
Skew:               -0.002    Prob(JB):                 5.19e-49
Kurtosis:           1.778    Cond. No.                  349.
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [ ]: from sklearn.linear_model import LogisticRegression # Import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix
import seaborn as sns
# Define the dependent and independent variables
X = df_cleaned[['Age', 'Gender', 'Review Rating']]
y = df_cleaned['Subscription Status']

# Fit the logistic regression model
model = LogisticRegression()
model.fit(X, y)

# Get the regression coefficients and intercept
coefficients = model.coef_
intercept = model.intercept_

# Print the regression coefficients and intercept
print('Regression coefficients:', coefficients)
print('Intercept:', intercept)

# Make predictions for specific values of the predictors
new_data = [[30, 1, 4], [45, 0, 3]]
predictions = model.predict(new_data)
print('Predicted subscription status:', predictions)
```

```

# Evaluate the model using accuracy, precision, recall, and F1-score
y_pred = model.predict(X)
accuracy = accuracy_score(y, y_pred)
precision = precision_score(y, y_pred)
recall = recall_score(y, y_pred)
f1_score = f1_score(y, y_pred)

print('Accuracy:', accuracy)
print('Precision:', precision)
print('Recall:', recall)
print('F1-score:', f1_score)

# Create a confusion matrix to visualize the model's performance
confusion_matrix = confusion_matrix(y, y_pred)
sns.heatmap(confusion_matrix, annot=True, fmt='d', cmap='Blues')
plt.title('Confusion Matrix for Subscription Status Prediction')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()

```

Regression coefficients:  $\begin{bmatrix} 1.42181364e-03 & 5.00078002e+00 & -1.95295013e-02 \end{bmatrix}$

Intercept:  $[-5.4207814]$

Predicted subscription status:  $[0 \ 0]$

Accuracy: 0.7309843400447428

Precision: 0.0

Recall: 0.0

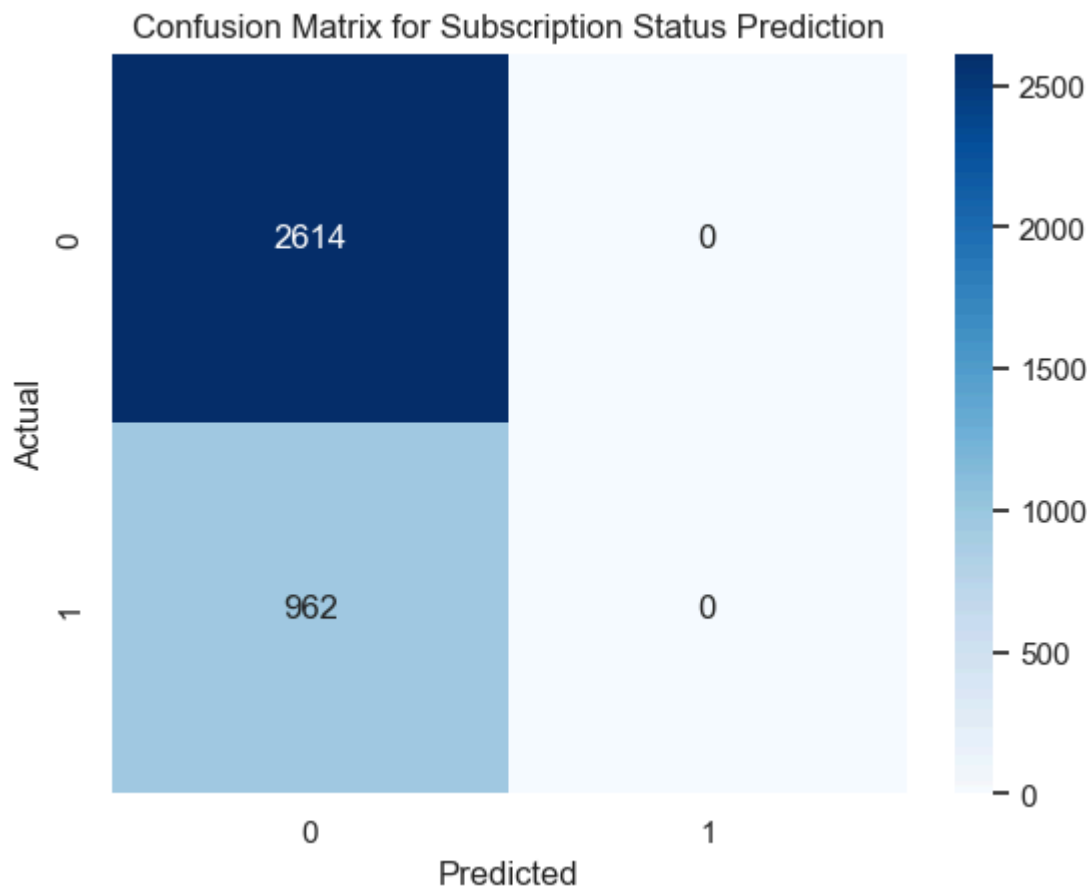
F1-score: 0.0

C:\Users\aravi\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.9\_qbz5n2kfra8p0\LocalCache\local-packages\Python39\site-packages\sklearn\base.py:493: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names

warnings.warn(

C:\Users\aravi\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.9\_qbz5n2kfra8p0\LocalCache\local-packages\Python39\site-packages\sklearn\metrics\\_classification.py:1497: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 due to no predicted samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, f"{metric.capitalize()} is", len(result))



```
In [ ]: # prompt: Decision Tree Regression
# • Fit a decision tree regression model to predict Purchase Amount (USD)
# based on Age, Gender, and Previous Purchases.
# • Visualize the decision tree.
# • Evaluate the model using metrics such as Mean Absolute Error (MAE), Mean
# Squared Error (MSE), and R-squared.
# • Compare the decision tree model with a multiple linear regression model.

import matplotlib.pyplot as plt
from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

# Define the dependent and independent variables
X = df_cleaned[['Age', 'Gender', 'Previous Purchases']]
y = df_cleaned['Purchase Amount (USD)']

# Split the data into training and test sets
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

# Fit the decision tree regression model
model = DecisionTreeRegressor(max_depth=3, random_state=42)
model.fit(X_train, y_train)

# Make predictions on the test set
y_pred = model.predict(X_test)

# Evaluate the model using MAE, MSE, and R-squared
mae = mean_absolute_error(y_test, y_pred)
```

```

mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print('Decision Tree Regression:')
print(f'MAE: {mae}')
print(f'MSE: {mse}')
print(f'R-squared: {r2}')

# Visualize the decision tree
from sklearn.tree import plot_tree
plt.figure(figsize=(10, 6))
plot_tree(model, feature_names=X.columns, filled=True, precision=2)
plt.title('Decision Tree for Purchase Amount Prediction')
plt.show()

# Compare the decision tree model with a multiple linear regression model
X_train_with_constant = sm.add_constant(X_train)
model_sm = sm.OLS(y_train, X_train_with_constant).fit()
y_pred_linear = model_sm.predict(sm.add_constant(X_test))

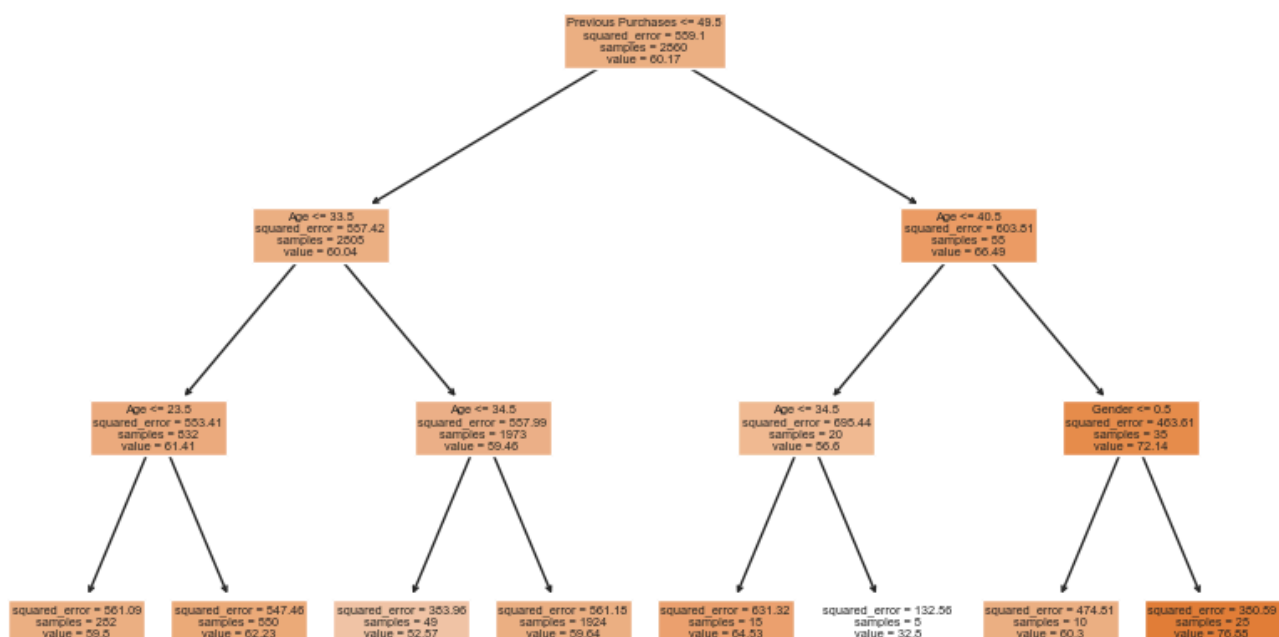
mae_linear = mean_absolute_error(y_test, y_pred_linear)
mse_linear = mean_squared_error(y_test, y_pred_linear)
r2_linear = r2_score(y_test, y_pred_linear)

print('Multiple Linear Regression:')
print(f'MAE: {mae_linear}')
print(f'MSE: {mse_linear}')
print(f'R-squared: {r2_linear}')

```

Decision Tree Regression:  
MAE: 20.244116169199966  
MSE: 551.6320392972503  
R-squared: -0.012119022902776555

Decision Tree for Purchase Amount Prediction



Multiple Linear Regression:

MAE: 20.24106704027441

MSE: 546.7865152388188

R-squared: -0.003228591009590165